

Reinforecment Learning: Notes and Selected Exercises

Somnath Sikdar

November 8, 2020

Contents

1	Markov Decision Processes	2
1.1	Introduction	2

Chapter 1

Markov Decision Processes

1.1 Introduction

Problems in reinforcement learning are typically formulated in the formal setting of a Markov Decision Process (MDP). An MDP is a five-tuple $\langle \mathcal{S}, \mathcal{A}, p, \mathcal{R}, \gamma \rangle$

- a set \mathcal{S} of states
- a set \mathcal{A} of actions
- a set \mathcal{R} of rewards
- a probability function $p: \mathcal{S} \times \mathcal{R} \times \mathcal{S} \times \mathcal{A} \rightarrow [0, 1]$
- a discount factor $\gamma \in [0, 1]$